

## Motivation for Hamilton Chapter 1


- **Describing the variability** and the observed distribution of data is the **required** first step of any data analysis.
- The **shape** of an univariate distribution can have **substantial impact** on the outcome of statistical procedures.  
E.g.: **Outliers** or **heavy tails** may detrimentally influence the outcome of model calibrations and estimations.
- Not accounting for the distribution of variables can force a researcher to redo their data analysis at a later state.
- **Most methods assume *symmetric* or preferably *normally* distributed variables.**
- Remember:
  - Data tell a story about the phenomena under investigation.
  - Always handle data and analysis results with a critical attitude and use common sense.
  - Always ask yourself: Do the data or the generated analysis results make sense?
- Transformations to symmetry are discussed in Chapter 1. Note, statisticians use many more transformations under in particular circumstances.  
E.g., we will encounter later the logit-transformation.

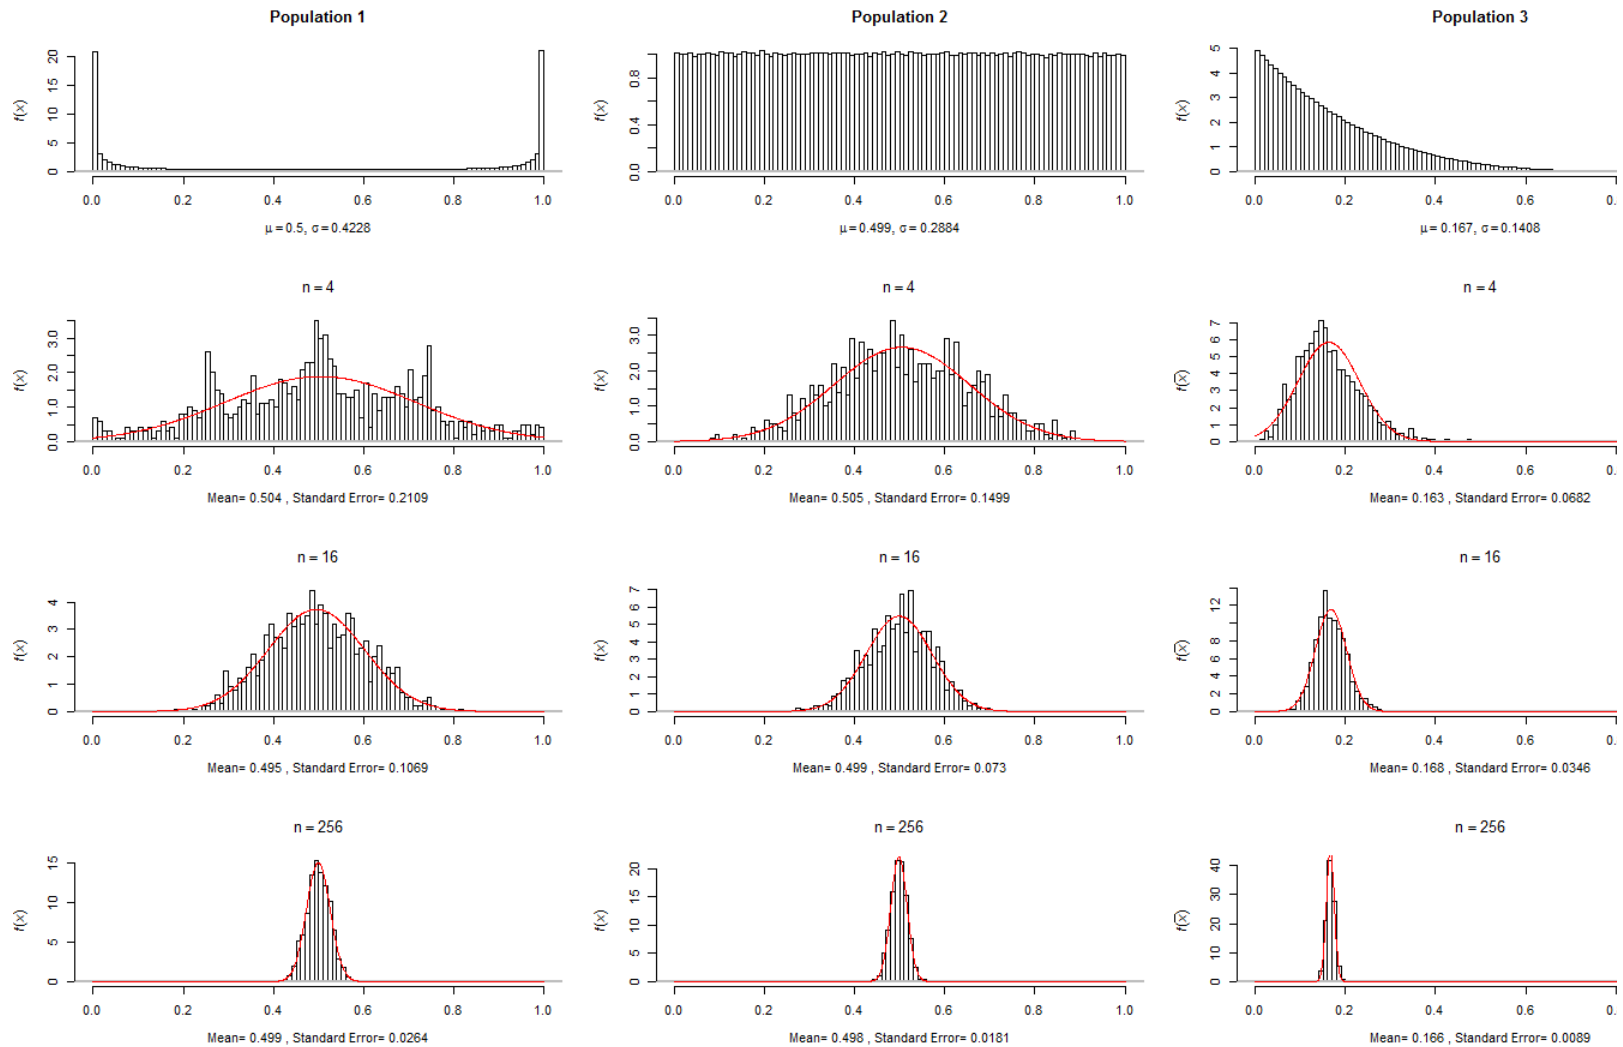
## Central Limit Theorem (skipped)

- Def. Central Limit Theorem: Let  $X_1, X_2, \dots, X_n$  be a **random independent** sample of size  $n$  drawn from an **arbitrarily distributed** population with **expectation  $\mu$**  and standard **deviation  $\sigma$** .

Then for large enough sample sizes  $n$ , the sampling distribution of  $\bar{X}$  is [a] asymptotically (i.e., as  $n \rightarrow \infty$ ) normal distributed [b] with  $\bar{X} \sim N(\mu, \sigma/\sqrt{n})$ .

Proof for independent sample objects: 
$$Var\left(\frac{1}{n} \cdot \sum_{i=1}^n X_i\right) = \frac{1}{n^2} \cdot \sum_{i=1}^n \underbrace{Var(X_i)}_{=\sigma^2} = \frac{1}{n^2} \cdot n \cdot \sigma^2 = \frac{\sigma^2}{n}$$

- Example: Central limit theorem with the -script **CENTRALLIMIT.R**:



## Review: The Shape of Distributions (skipped)

- Distributions can be distinguished with regards the **balance** of their left and right tails:
  - **Symmetric** distributions. Tails are balanced into either direction from a central value.
  - **Negatively skewed** distributions (long tail into the negative direction)
  - **Positively skewed** distributions (long tail into the positive direction). These distributions frequently emerge for variables with a binding lower origin (like zero income).
  - Extreme skewness may hint at **outliers** that do not match the rest of the observed data.
- The number of meaningful clusters of observations is described by the term **modality**:
  - **Uni-modality** refers to just one peak
  - **Bi-modality** refers to two outstanding peaks
  - **Multimodality** refers to more than two outstanding peaks.
- **Multimodality** may hint at a **heterogeneous** underlying data generating process.

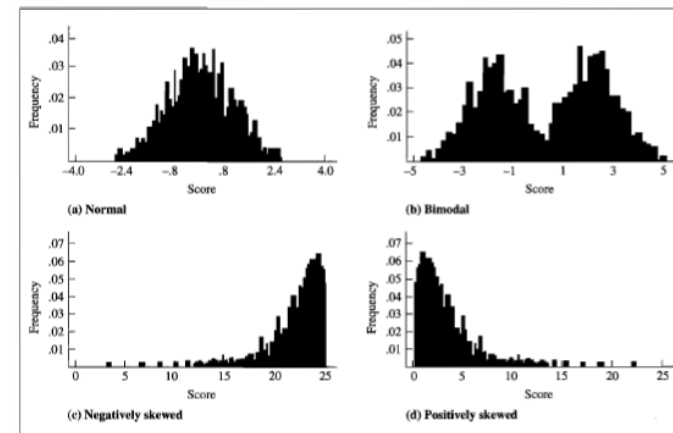


Figure 3.9  
Shapes of frequency distributions: (a) Normal; (b) Bimodal; (c) Negatively skewed; (d) Positively skewed

## Quantiles and Percentiles (skipped)

- Technically, quantiles and percentiles are generated from a **sorted list** of the original data points  $x_{[1]} \leq x_{[2]} \leq x_{[3]} \leq \dots \leq x_{[n-1]} \leq x_{[n]}$  where each observations has an assigned rank  $i \in \{1, 2, \dots, n\}$ , with  $i = 1$  for the smallest observation and  $i = n$  for the largest observation.
- For a give data value  $x_{[i]}$  the **percentile** approximates the proportion of sample observations less or equal to  $x_{[i]}$ , that is,  $p_{[i]} = \frac{i - \frac{1}{2}}{n} \approx \Pr(X \leq x_{[i]}) = \int_0^{x_{[i]}} f(x) \cdot dx$ .
- A **quantile** is that observed data value of a distribution, which is associated with a particular percentile point.
- Important quantiles are:
  - 0.25 quantile also called  $Q_1$  quartile (25 % of the observations are smaller or equal to this quantile value)
  - 0.50 quantile also called the median (50 % of the observations are smaller or larger than the given quantile value)
  - 0.75 quantile also called  $Q_3$  quartile (75 % of the observations are smaller or equal to this quantile value and 25 % of the observations are larger than this value)
  - A measure of spread is the inter-quartile range:  $IQR = Q_3 - Q_1$

## Box-Plots (skipped)


- Construction of the box-plot
  - Draw a **box** from  $Q_1$  to  $Q_3$ . Mark the **median**  $Q_2$  in the center of the box with a line.
  - Definition of **adjacent values**  $x_{low}^{adj} = \min(x_{[i]} \in (Q_1, Q_1 - 1.5 \cdot IQR) \text{ plus } x_{[i]} \text{ in dataset})$  and  $x_{high}^{adj} = \max(x_{[i]} \in (Q_3, Q_3 + 1.5 \cdot IQR) \text{ plus } x_{[i]} \text{ in dataset})$ .

The term  $x \in (a, b)$  means, all  $x$ -values in the interval between  $a$  and  $b$ .

Draw the “fences” so they just include the smallest and largest data values  $x_{low}^{adj}$  and  $x_{high}^{adj}$ , respectively.

- **Outliers** are in the interval  $[1.5 \cdot IQR, 3.0 \cdot IQR]$  starting from  $Q_1$  below or  $Q_3$  above, respectively.  
**Severe outliers** are beyond that range ( $> 3.0 \cdot IQR$ )

- Use of box-plots:
  - Easy visual description of the distribution of a variable and potential outliers



- Comparison of distributions for several variables side-by-side.

## Quantile-Normal Plot

- Calculate the theoretical quantiles of a normally distributed random variable  $X_{[i]}$  (assuming the mean  $\mu$  and the variance  $\sigma^2$  were estimated from the sample data) based on the given percentiles  $p_{[i]}$  of the observed variable  $Y_{[i]}$ .
- Quantile-Normal Plot:** Plot the theoretical normal distribution quantiles  $X_{[i]}$  on the abscissa (X-axis) against their matching empirical distribution of  $Y_{[i]}$  on the ordinate (Y-axis).

Interpretation:

- Diagonal with slope 1 => equal distributions;
- Not a straight-line => different shapes

## Properties of Arithmetic Mean (skipped)

- Implications of the **zero-sum** property  

$$\sum_{i=1}^n (Y_i - \bar{Y}) = 0$$
 Assuming the mean is known, then  $n-1$  observation can vary freely, whereas we can predict the last observation with certainty.

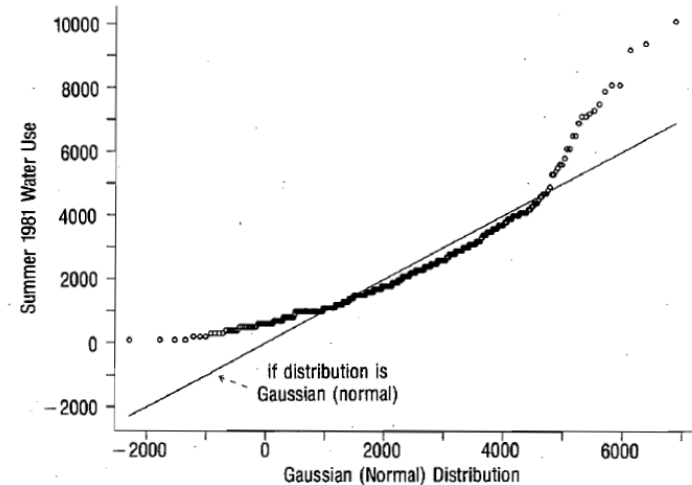


Figure 1.9 Quantile-normal plot of household water use (positively skewed).

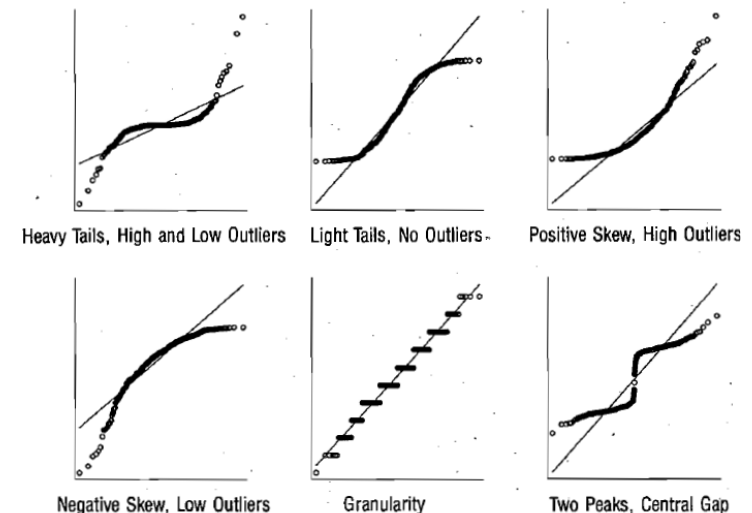


Figure 1.10 Quantile-normal plots reflect distribution shape.

$$\sum_{i=1}^n (Y_i - \bar{Y}) = 0$$

$$\Rightarrow \sum_{i=1}^n Y_i = n \cdot \bar{Y}$$

$$\Rightarrow Y_n = n \cdot \bar{Y} - \sum_{i=1}^{n-1} Y_i$$

That implies that we loose on degree of freedom.


- Implication of the **least squares property**  $\min_{\theta} \sum_{i=1}^n (Y_i - \theta)^2 \Rightarrow \theta = \bar{Y}$ .

Large deviations have a strong impact on the estimated mean, variance etc. because the large deviations are squared

$\Rightarrow$  Thus, large deviations pull the mean into their direction.

$\Rightarrow$  Standard deviations are drastically inflated.

- Lacking any other information, the arithmetic mean will become best **predictor** for the variable under question.
- The deviations from the mean are the **unexplained** part or the **residuals** of the observations, i.e.,  $y_i = \bar{y} + \varepsilon_i$ .



- Definition of **total sum of squares**:  $TSS = \sum_{i=1}^n (Y_i - \bar{Y})^2$  or  $TSS = \sum_{i=1}^n Y_i^2 - n \cdot \bar{Y}^2$ .

variance



- Why is the population variance estimated  $(n - 1)$  in the denominator, that is, by  $s^2 = TSS/(n - 1)$  :

**Explanation 1:** If we **calculate the mean** from the sample then there are only  $n - 1$  "degrees of freedom" left because of the **zero sum property** of the mean.

**Explanation 2:** **The mean is calculated to minimize the TSS.**

Thus the sample mean always fits the observed sample data better than any **unobserved but true** population expectation  $\mu$ .

For the true expectation  $\mu$ , the TSS would be slightly larger. That is why the sample TSS needs to be inflated by dividing it by a slight smaller value than  $n$ , that is,  $n - 1$ .

- **Standard deviation** measures the variation in **original units** rather than in squared units.

## Review: Skewness (skipped)

- Why does the distribution of the water consumption in the Concord dataset deviate from the normal distribution?

Reason: Fixed lower bound (negative consumption impossible).

- **Skewness** and bounded/truncated distributions: For skewed distributions the notion of the center of the distribution (mean) becomes ambiguous and the **median** may be a better representation of the central tendency in the data.

- The **skewness** is defined by  $skew(X) \equiv \frac{\sum_{i=1}^n (x_i - \bar{x})^3}{n \cdot s_X^3}$

- The **normal distribution has a skewness of 0.**

## Box-Cox Transformation

- This lecture focuses on the more general *Box-Cox* transformation rather than the slightly simpler *power*-transformation, which is discussed in Hamilton.  
For both transformations the general interpretation of the parameter  $\lambda$  does not change.
- **Causes for *extreme observations*:** [a] **skewed distributions**, [b] measurement or recoding errors, [c] extreme but feasible events (perhaps not belonging to the population under investigation).
- The *power*-transformation presented in book and the **Box-Cox transformation only work for  $X$ 's larger than zero**.
- Thus the standard Box-Cox transformation will not work for data values  $X$  that are zero or negative.
  - In order to avoid this problem, **a constant such as  $\min(X)$  or, say, 5% quantile, needs to be added to  $X$  with negative values, in order to make it positive.**
  - However, if the constant is too small leading to positive but close to zero values, outliers may be introduced.
  - On the other hand, choosing the constant too large, may make the transformation to normality ineffective.
- The Box-Cox transformation is a generalization of the power transformation:  $Y = \frac{X^\lambda - 1}{\lambda}$  and for  $\lambda = 0$  we get  $y = \ln(x)$ . See also note 11 on page 28 in Hamilton
- $\lambda > 1$  reduce negative skewness, whereas  **$\lambda < 1$  reduce positive skewness**.  
Remember: Positive skewness is very common for variables with a natural bound of zero.

- If power  $\lambda < 0$  then all values are multiplied by a negative number to preserve the order of observations.

This explains the value  $\lambda$  in the denominator of the Box-Cox transformation

FOX Fig 4.1




Figure 4.1. The family of power transformations  $X'$  of  $X$ . The curve labeled  $p$  is the transformation  $X^{(p)}$ , that is,  $(X^p - 1)/p$ ;  $X^{(0)}$  is  $\log_e X$ .

- Note: R's function `car::powerTransform()` is performing several statistical tests whether a variable either needs to be transformed or whether a *log*-transformation is sufficient by using the likelihood ratio test (LR) principle:
  - The first LR tests the null hypotheses  $H_0: \lambda^{optimal} = 0$ . If we cannot reject the null hypothesis then we should tentatively **work with a log-transformation** to achieve normality/symmetry.
  - The second LR tests the null hypotheses  $H_0: \lambda^{optimal} = 1$ . If we cannot reject the null hypothesis then we should tentatively **should work with an untransformed variable** because it is approximately symmetric.

- The Wald confidence interval provides the 95% probability range within which the true population transformation parameter  $\lambda$  lies.

## Handling Transformations with Negative Data Values

- In case some data values are negative or zero a small constant  $\gamma$  can be added to the data  $X$  (see the `?car::bcPower( )` and Fox & Weissberg pp 161-162 for the **bcnPower** transformation family)
- A more informed way avoiding some of the problems by just adding a constant is to first transform the data by:

$$z(X, \gamma) = \frac{(X + \sqrt{X^2 + \gamma^2})}{2} \text{ with}$$

- The transformation  $z(X, \gamma)$  is monotonic (i.e., if  $x_1 < x_2$ . then  $z(x_1, \gamma) < z(x_2, \gamma)$ )
- For large positive  $X$  relative to  $\gamma$  ( $X \gg \gamma$ ) the transformation is approximately linear with  $z(X, \gamma) \approx X$ .
- If  $\gamma = 0$  then  $z(X, \gamma) = X$  for  $X > 0$  and  $z(X, \gamma) = 0$  for  $X \leq 0$ .
- Subsequently, once the  $\gamma$ -parameter is determined a standard Box-Cox transformation is applied to  $z(X, \gamma)$ .

## LOESS Smoother of $Y \sim X$ Relationships

- Many of R's scatterplot functions not only show a linear regression fit through the data cloud but also show a locally smoothed loess-curve:
  - In essence the sliding window moves over the value range of X.
  - In each window a local regression line is estimated.
  - These local regression lines are “splined” together into the smooth loess curve over the whole value range of X




FIGURE 9.17  
MA-plot with curve obtained with loess.




FIGURE 9.16  
Illustration of how loess estimates a curve. Showing 12 steps of the process.